

Mapbox Vector Tile Specification 2.0

Blake Thompson - Software Engineer, Mapbox



About Me

- Developer at OKC Mapbox Office
 - Mapnik
 - Node Mapnik
 - Mapnik Vector Tile
- Author of Mapbox Vector Tile Specification



Coffee during morning scrum

What are Vector Tiles

Binary format for encoding geometries and their metadata

- Points
- Lines
- Polygons

Structure of Vector Tiles

- Layer
 - Feature
 - Geometry

Lots of Companies Using Vector Tiles!



What I Will Cover

- Why use Vector Tiles?
- Why we made Version 2.0 of the Spec.
- What changed in 2.0?
- Common Misconceptions about Vector Tile Spec.
- Future of the Spec.

Why use Vector Tiles?

(besides all the cool features)

Your #1 Limiting Resource is
\$\$\$\$\$\$\$\$\$\$



Cost of Making a Map

Steps To Make Map:

- Query
- Render
- Transfer

COST FORMULA:

R = Total Number of Requests (Work)

q = query

r = render

t = transfer

s = storage of data

$$\text{\$ (Cost)} = R (q + r + t + s)$$

WMS

- R (Requests) increases a lot if the map is interactive
- Query per request
- Rendered on demand
- Transfer is an image (t is large)
- **Storage must scale with R because databases must scale with requests**

Cost is linear but STEEP!

$$\text{\$} = R (q + r + t + s)$$





Image Tile Rendering from Database

- Requests are reduced because caching can be done!
 - R_c = Requests with Caching
- Transfer is an image (t is large)
- **Storage must scale with R because databases must scale with requests**

Cost is linear and steep but caching can greatly reduce costs! Cheaper than WMS!

$$\text{\$} = R_c (q + r + s) + R (t)$$




Serving Static Image Tiles

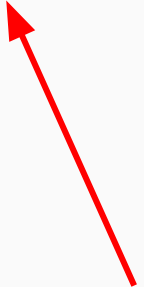
- Create all the images before serving!
 - S = Number of styles
- **Storage can become quite massive as you are storing both the source data and the image tiles!**
- Great for static datasets with high traffic
- Query and Rendering are now based on number of styles!

Cost is linear and almost constant if you have a very small number of styles and static data. However, initial cost is HIGHER.

$$\text{\$} = R(t) + S(s + q + r)$$



Super Easy to Scale!




Extremely expensive if you need a lot of different map styles

Vector Tile Server Side Rendering

- Precut all the geometry into Vector Tiles and store!
- Storage (s) is smaller as VT take much less space than images
- Storage not repeated per style!
- Great for large number of different styles!

Cost increases based on requests, but much cheaper if you have a lot of styles of the same data!

$$\text{\$} = R_c (r) + R (t) + (q + s)$$



Much smaller storage cost then storing images

Vector Tile Client Side (GL) Rendering

- Transfer size is smaller as you are no longer sending images!
- **No rendering cost as client's hardware does the rendering!**

$$\$ = R(t) + (q + s)$$

Very low cost!!!

Vector Tiles on the Client

- Query data on map
- Interactive maps
- Smooth Transition (no transition between zoom levels of data)
- Easy to still render tiles on the server side
- Very quick transfer of data to client!

Why Vector Tile 2.0?

Make it a “Real” Specification

- Version 1.0 left many things unexplained
- Implementations were simply based off Mapnik Vector Tile library

Vector Tile Specification 1.0 fits in one slide

Vector Tiles use [Google Protocol buffers](#) as a container format. It is exclusively geared towards square pixel tiles in [Spherical Mercator projection](#).

A vector tile can consist of one or more named layers and containing one or more features.

Features contain an id, attributes, and geometries: either point, linestring, or polygon.

Geometries are stored as an a single array of integers that represent an command,x,y stream (where command is a rendering command like move_to or line_to). Commands are encoded only when they change.

Geometries are clipped, reprojected into spherical mercator, converted to screen coordinates, and [delta](#) and [zigzag](#) encoded.

Feature attributes are encoded as key:value pairs which are dictionary encoded at the layer level for compact storage of any repeated keys or values. Values use [variant](#) type encoding supporting both unicode strings, boolean values, and various integer and floating point types.

```
// Contains a stream of commands and parameters (vertices). The
// repeat count is shifted to the left by 3 bits. This means
// that the command has 3 bits (0-7). The repeat count
// indicates how often this command is to be repeated. Defined
// commands are:
// - MoveTo:    1    (2 parameters follow)
// - LineTo:    2    (2 parameters follow)
// - ClosePath: 7    (no parameters follow)
//
// Commands are encoded as uint32 varints. Vertex parameters
// are encoded as deltas to the previous position and, as they
// may be negative, are further "zigzag" encoded as unsigned
// 32-bit ints:
//
//   n = (n << 1) ^ (n >> 31)
//
// Ex.: MoveTo(3, 6), LineTo(8, 12), LineTo(20, 34), ClosePath
// Encoded as: [ 9 6 12 18 10 12 24 44 15 ]
//           |   |   |           `> [00001 111] command type 7 (ClosePath), length 1
//           |   |   |           ===== relative LineTo(+12, +22) == LineTo(20, 34)
//           |   |   |           ===== relative LineTo(+5, +6) == LineTo(8, 12)
//           |   |   |           `> [00010 010] = command type 2 (LineTo), length 2
//           |   |   |           ===== relative MoveTo(+3, +6)
//           |   |   |           `> [00001 001] = command type 1 (MoveTo), length 1
//
// The original position is (0,0).
repeated uint32 geometry = 4 [ packed = true ];
```

Community Direction

- Vector Tiles should be useable across different vendors
- Community should participate in the future of the specification

Friendlier Data

- Invalid geometry data complicates display and analysis
 - Self Intersections
 - Self tangency
 - Repeated points
- No repeated layer names
- No repeating feature IDs

A Process for Changing Specification

- An understood process for changing the specification
- Definitive way to Version specification

What Changed in 2.0?

Compatibility

- Version 1.0 decoders will be able to decode 2.0 tiles
- Version 1.0 tiles should be able to be converted to 2.0 tiles
 - Correct Geometry
 - Remove “Empty” Layers and Features

Identify Vector Tiles

- File Extension is “.mvt”
 - “.vector.pbf” is not descriptive
 - Protobuffers requires a knowledge of data structure
- MIME Type
 - application/vnd.mapbox-vector-tile
 - Registered with IANA

Projections

- Each Vector Tile has its own “projection”
 - Data in a coordinate system relative to the top left corner of the vector tile.
- Clarified that data of almost **any** projection could be put into a vector tile
 - Most data still utilizes mercator!

Multi Type Geometries

- Well Defined Multi Types
 - Multi Point
 - Multi Linestring
 - Multi Polygon
- No “Geometry Collections”


Geometry Encoding

- Completely defined all “commands” and how to encode geometry in Vector Tiles
- Provided many examples in specification of encoding geometry
- A “LineTo” command must make a line
 - Can not result in a repeated point!



Winding Order

- Polygons now have a clearly defined winding order
- Area Positive is Exterior Ring
- Area Negative is Interior Ring
- Polygon
 - One Exterior Ring
 - Followed By Interior Rings
- New Polygon starts with each Exterior Ring


Polygon with Hole



Invalid Polygon



MultiPolygon



Polygon Validity

- Intersections have been a common problem
- Cause extensive problems for rendering applications

V 2.X Requires:

- No Self Intersections
- No Self Tangencies
- No Repeated Points
- "OGC Valid"


1 / 4

Polygon with a "hole"



Benefits of Valid Geometries

- Improved GL rendering
 - Tessellation errors removed
 - Faster processing
- Analysis Tools using VTs
 - More Accurate
 - No clean up required



Example Tessellation Error in GL Client from Invalid Geometry

Common Misconceptions

Collection of Tiles

- Projection or location of a tile is **NOT** encoded in a tile
 - Typically determined by filename or data when stored in database.
- The specification does not define how collection of tiles should be formed
 - This might be formed later in its own specification?
 - Just like PNG format does not specify for image tiles

Vector Tiles as “Source Data”

- Vector tiles are not intended to be **precise**
 - Lose precision due to rounding to integer coordinates of Vector Tile coordinates
 - Optimized for size, speed, and display
- Best viewed as static datasource
 - Vector tiles are not designed to be edited.

Compression

- Vector Tiles are encoded efficiently but are **not** required to be compressed
 - Compression is not part of the specification!
- Gzip compression is common
 - Extension should be “.mvt.gz”
 - Web servers often send vector tiles compressed!
 - Node Mapnik automatically recognizes gzip compression

Future of the Specification?

(Experiments taking place - no promises)

Raster/Images in Vector Tiles

- Serve raster data as a layer with geometry
 - Grayscale rasters
 - Encoding of rasters?
- Subject of Experiences
- Might be best left out of specification

Other Possible Additions

- Feature Collections
- Null Values
- Improved Compression of Geometry

Questions?



Backup Slides

Vector Tiles vs PNG

- Store Geometry
- Not Human Readable
- **Not ready for display**
- No Projection
- Not Precise

- Store Raster (Image)
- Not Human Readable
- **Ready for display**
- No Projection
- Not Precise

Vector Tiles vs GeoJSON

- Store Geometry
- Not Human Readable
- Not ready for display
- No Projection
- Not Precise
- **Very Small (< 1KB)**

- Store Geometry
- **Human Readable**
- Not ready for display
- **Can Provide Projection**
- **Precise**
- Can be very large (> 1 MB)